

SIA-TP5

Deep Learning

Grupo 6

- Francisco Sendot
- Lucia Dígon
- Martín E. Zahnd
- Juan Ignacio Fernández Dinardo

01

Autoencoder



Lista de caracteres “font.h”

h e l l o w o r l d

p a n d a z u m i x e r

Búsqueda de arquitectura: variando



Capas

Probamos distintas cantidades de capas de distintos tamaños



Optimizadores

Probamos con Adam y Gradient Descent



Epochs

Variamos la cantidad de epochs



Func. activación

Tanh, sigmoid, logistic



Learning rate

Cambiamos el learning rate



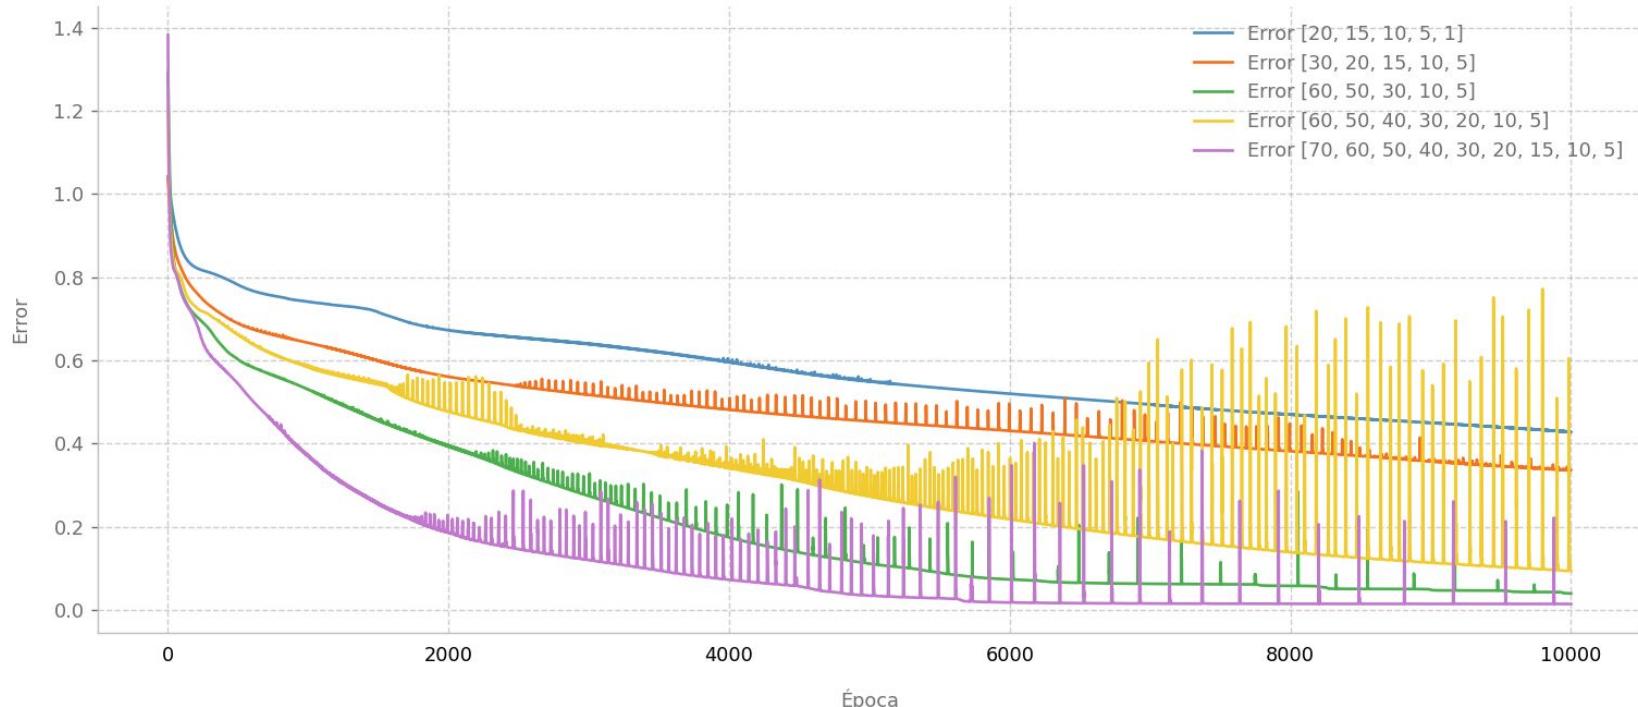
Entrenamiento

Batch y mini batch

Variando arquitectura (capas)

Error durante el entrenamiento del Autoencoder

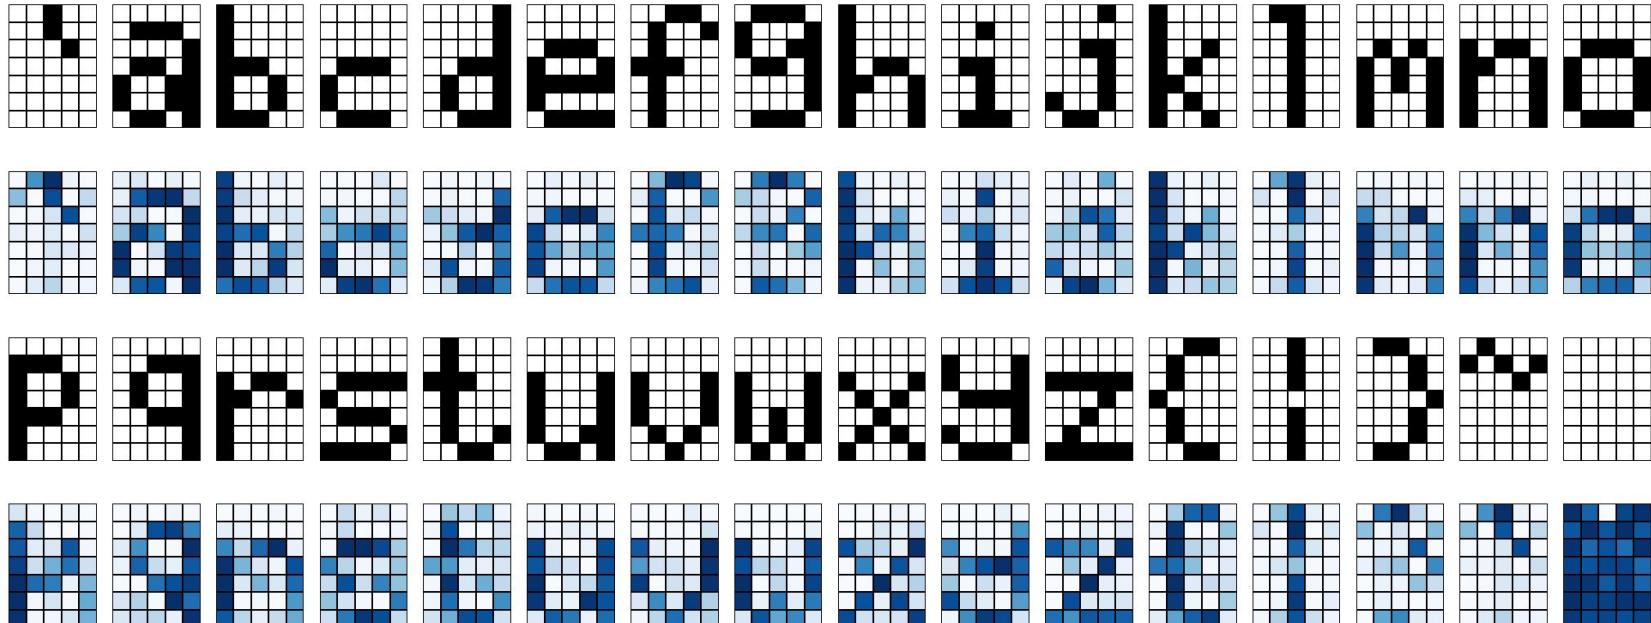
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000



Variando learning rate: 0.001

Input vs Autoencoder Output Heatmaps

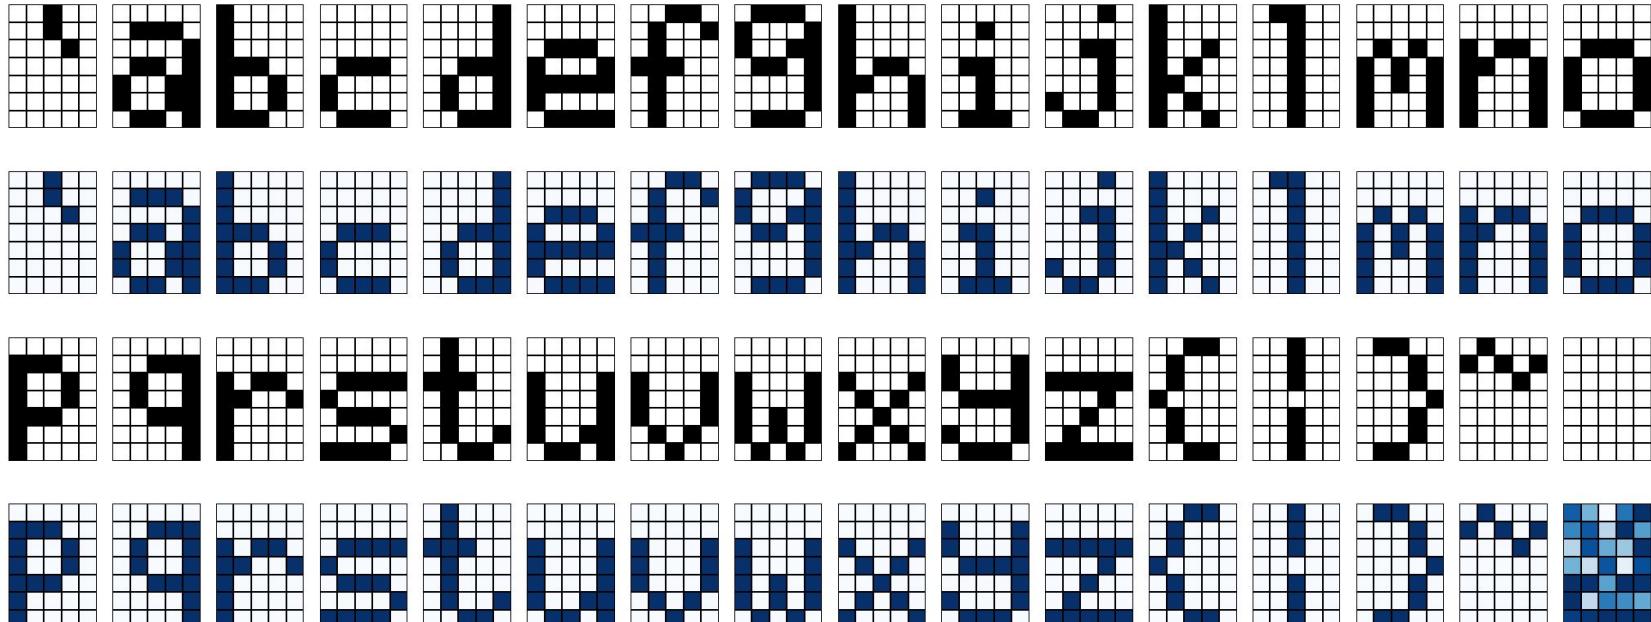
Learning Rate = 0.001 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



Variando learning rate: 0.005

Input vs Autoencoder Output Heatmaps

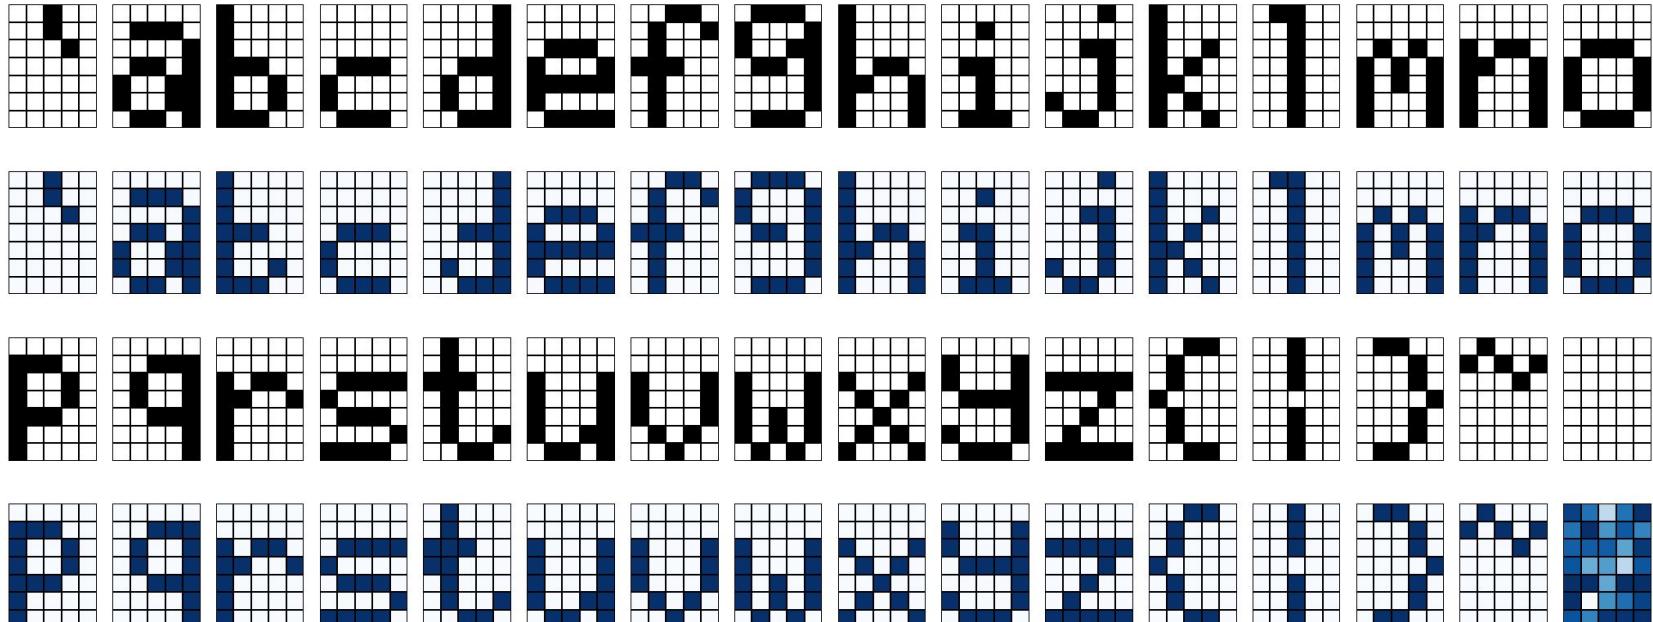
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



Variando learning rate: 0.008

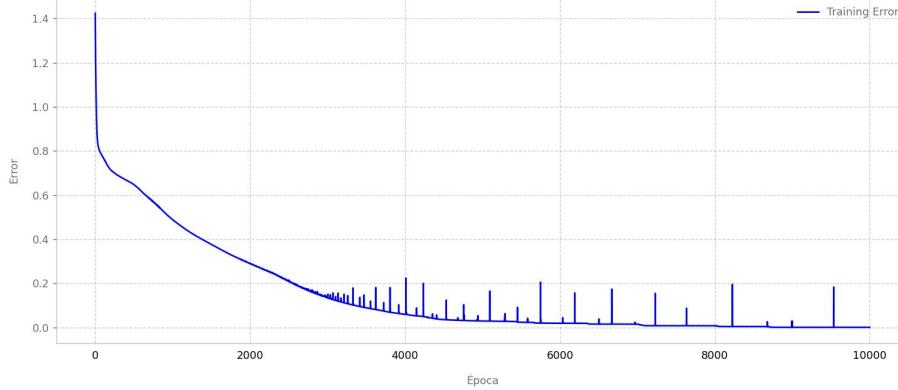
Input vs Autoencoder Output Heatmaps

Learning Rate = 0.008 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



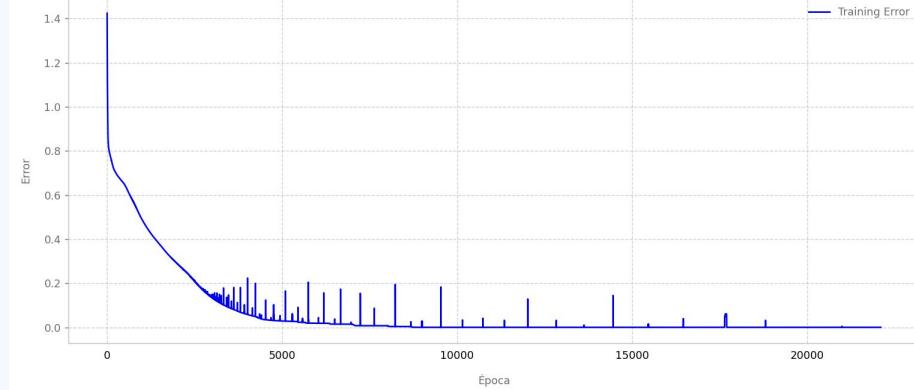
Variando Epochs

Error durante el entrenamiento del Autoencoder
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000



10.000

Error durante el entrenamiento del Autoencoder
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 30000

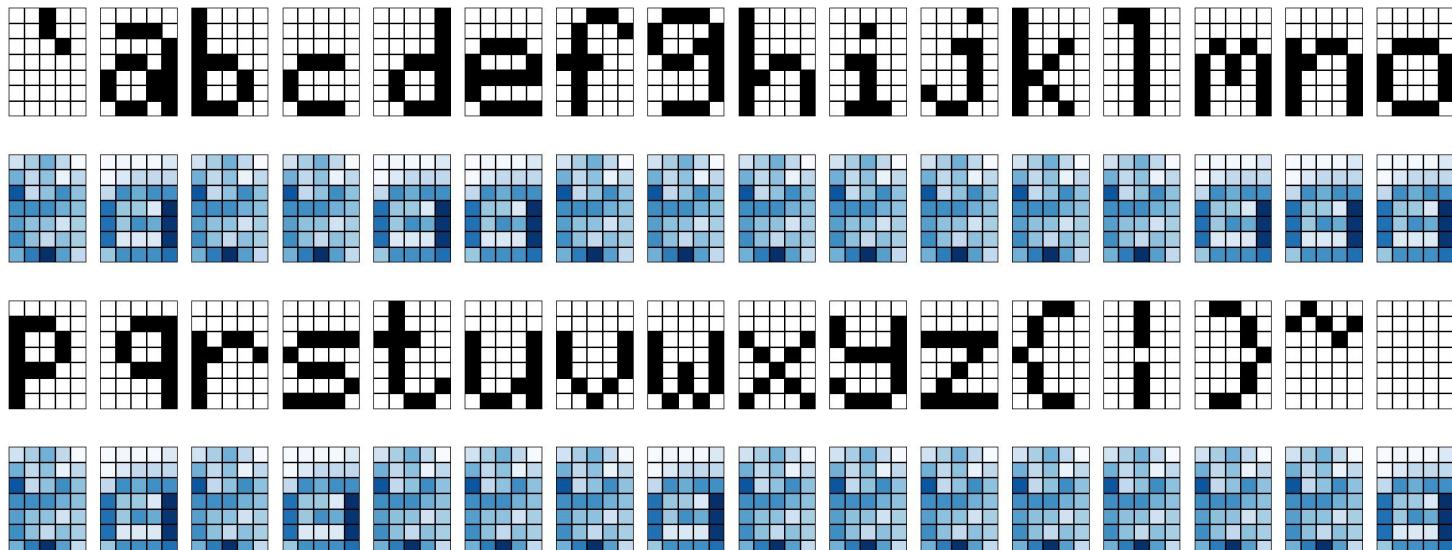


30.000

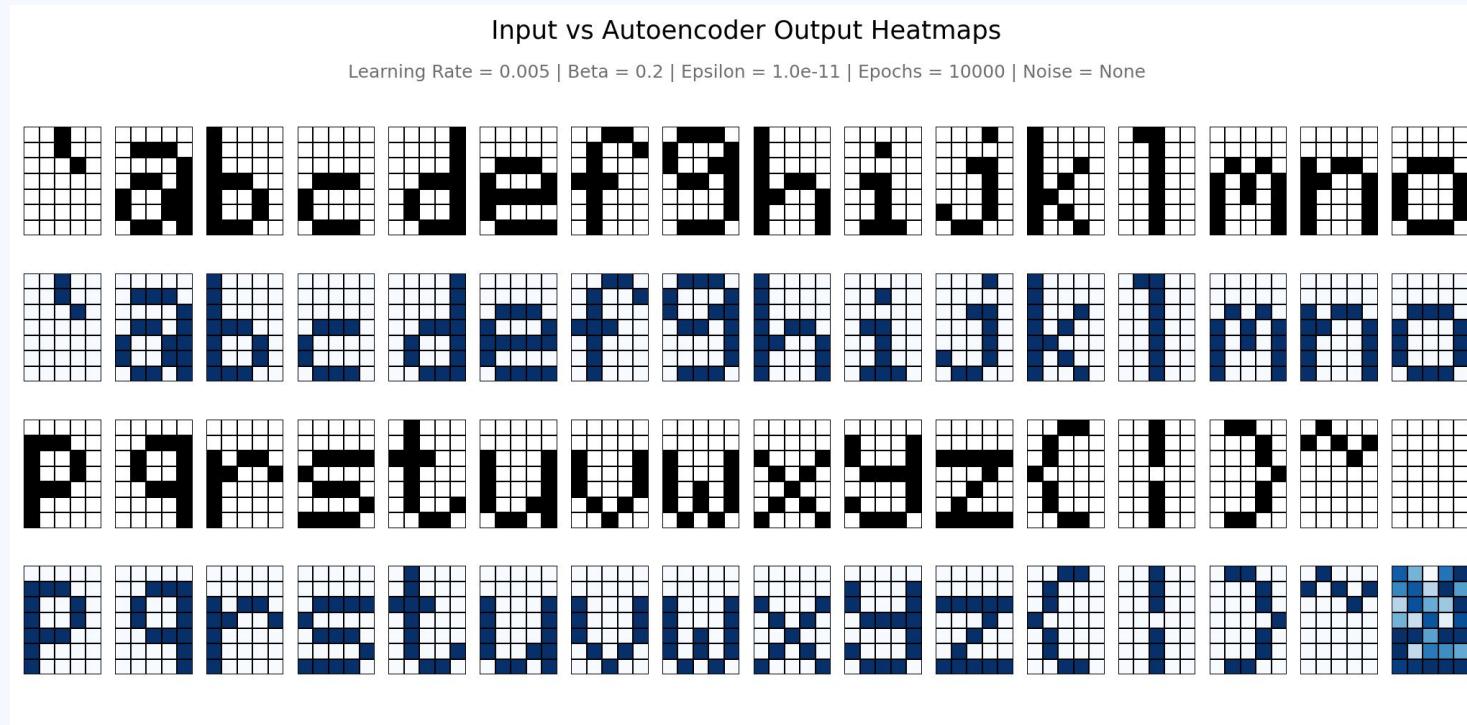
Variando Optimizadores: Gradient Descent

Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



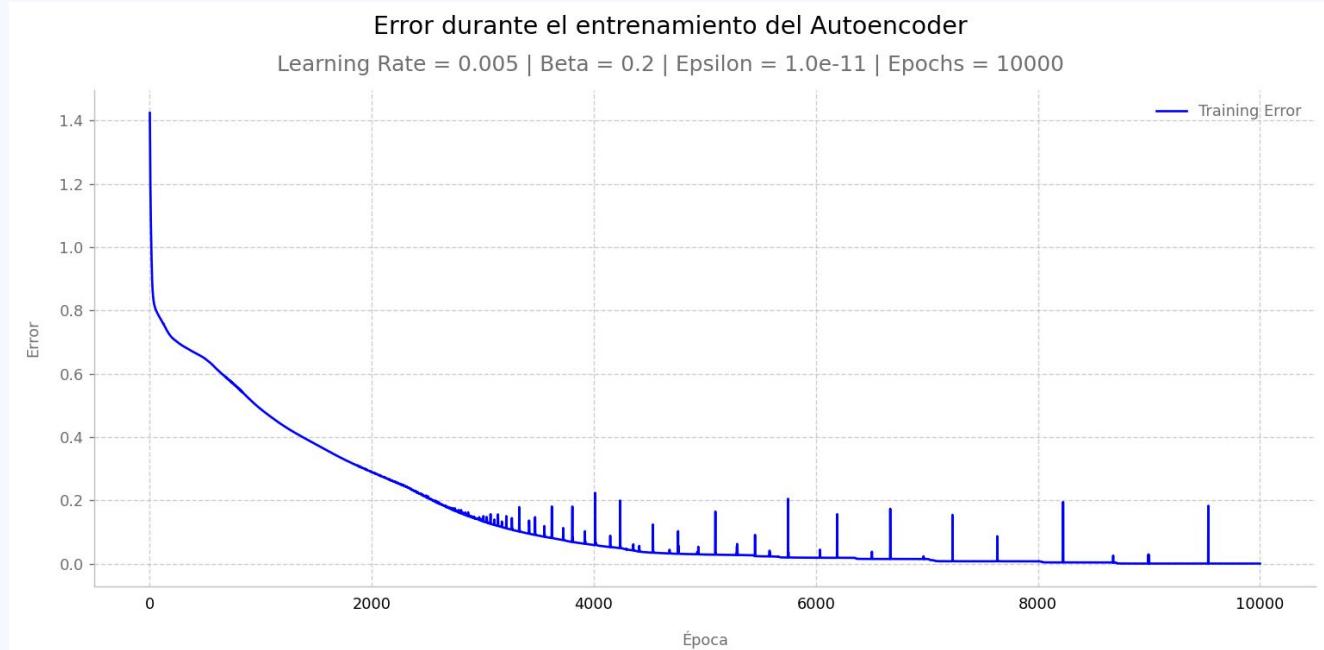
Variando Optimizadores: Adam



Arquitectura elegida

35-60-50-30-10-5-2-5-10-30-50-60-35

- **Epochs:** 10.000
- Tanh
- Adam
- **Learning rate:** 0,005

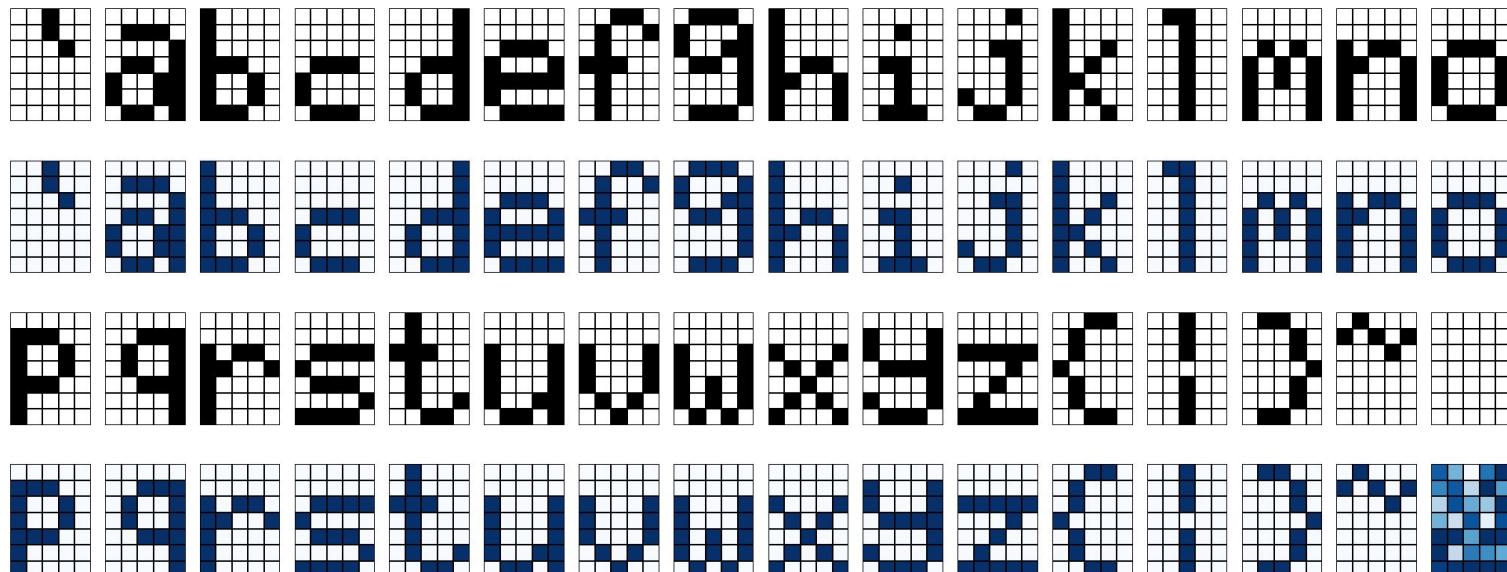


Arquitectura elegida

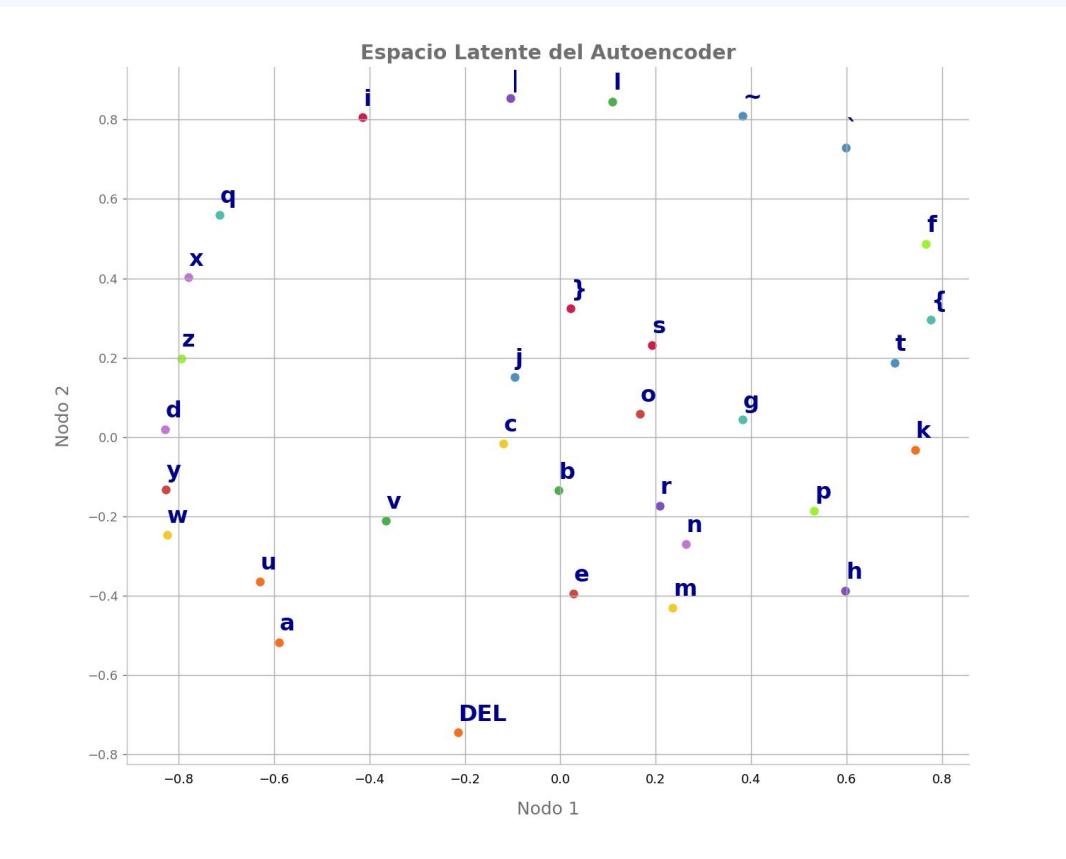
35-60-50-30-10-5-2-5-10-30-50-60-35

Input vs Autoencoder Output Heatmaps

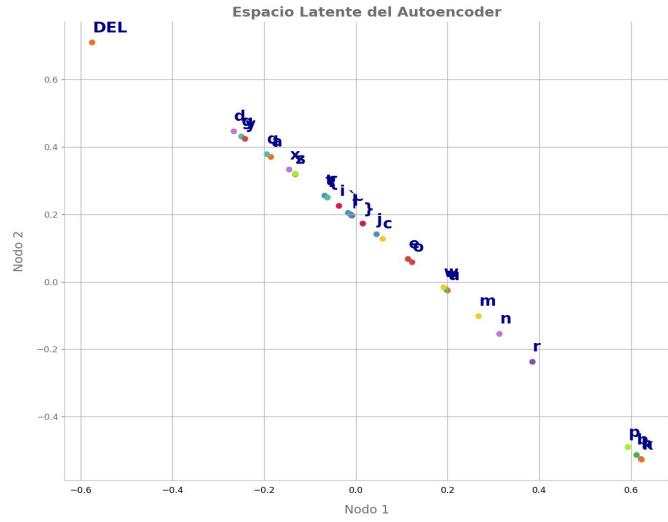
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



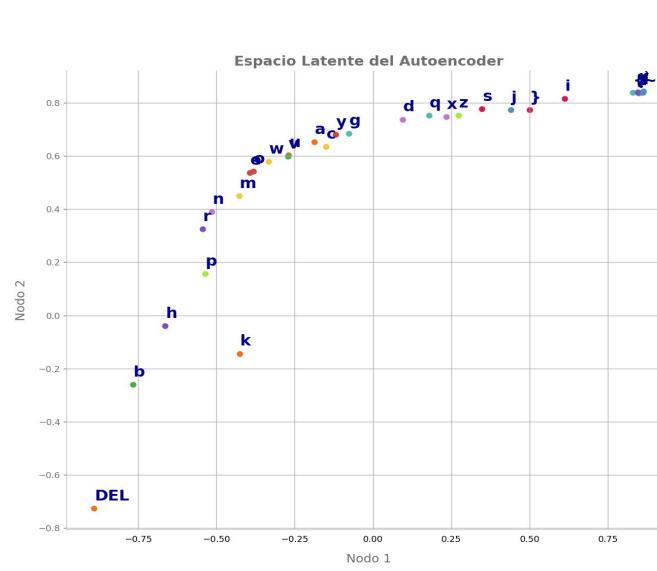
Datos de entrada en el espacio latente



Espacios latentes para arq que no aprendieron

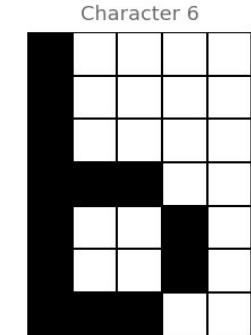
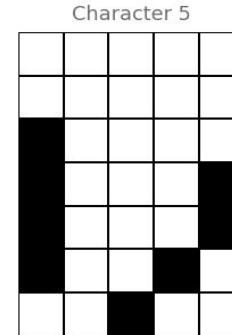
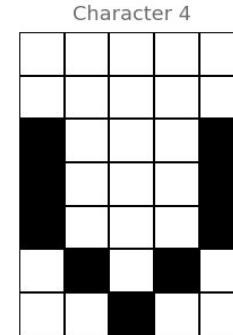
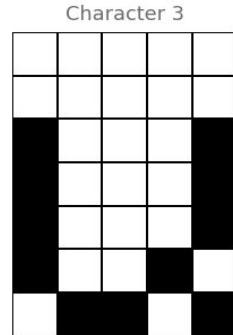
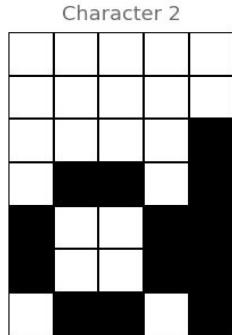
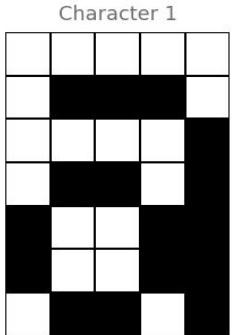


35-20-15-10-5-1-2

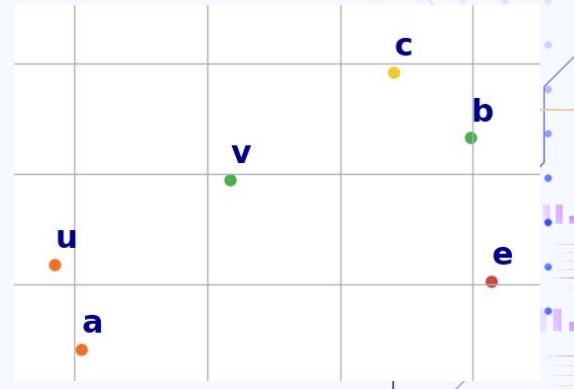


35-30-20-15-10-5-2

Generación nueva letra



- Se obtiene el la codificación dentro del espacio latente de dos caracteres, en este caso a y b
- Se toma un “paso”, y se recorre el espacio latente desde uno a otro.



Conclusiones

- Para diferentes arquitecturas, varía el espacio latente, más si no llega a aprender bien.
- En el espacio latente, se puede ver las letras que tienen codificaciones similares
- Llega un punto que seguir bajando el learning rate hace que el autoencoder no llegue a aprender bien

02

Denoising

Autoencoder



Ruido

- Utilizamos ruido Gaussiano
- Tomamos cada pixel prendido y alteramos los de su entorno siguiendo

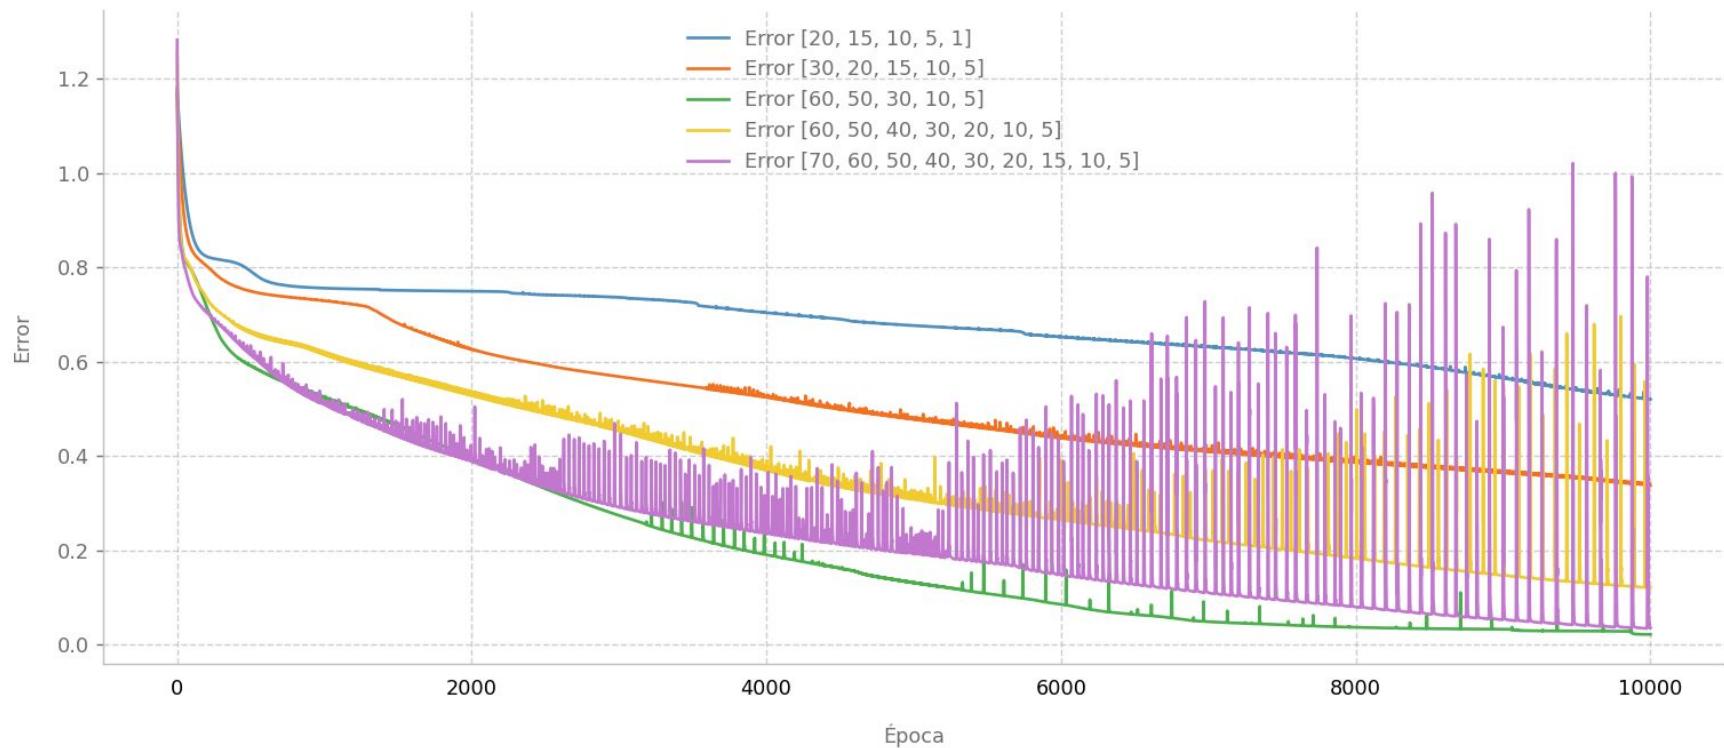
$$\mathcal{N}(\mu, \sigma^2) * \text{intensity} / 2$$

Con $\mu = 1$ y $\sigma^2 = 1$

Ruido - Arquitectura

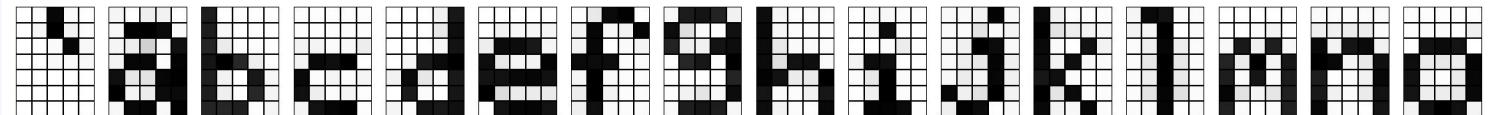
Error durante el entrenamiento del Autoencoder

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000

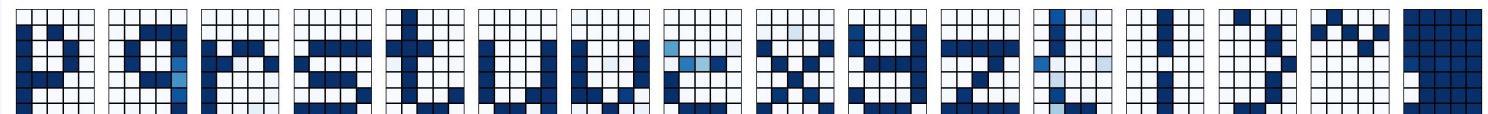
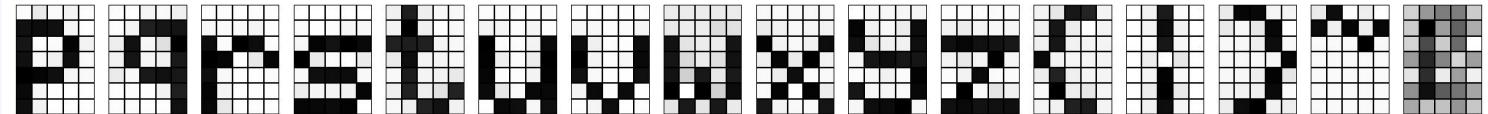
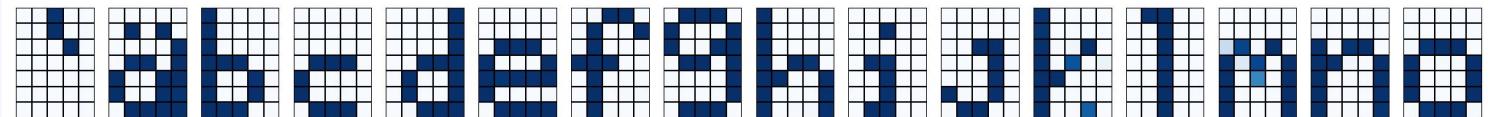


Dataset - Ruido 0.1

Entrada con ruido
del entrenamiento



Aprendizaje



Input vs Autoencoder Output Heatmaps

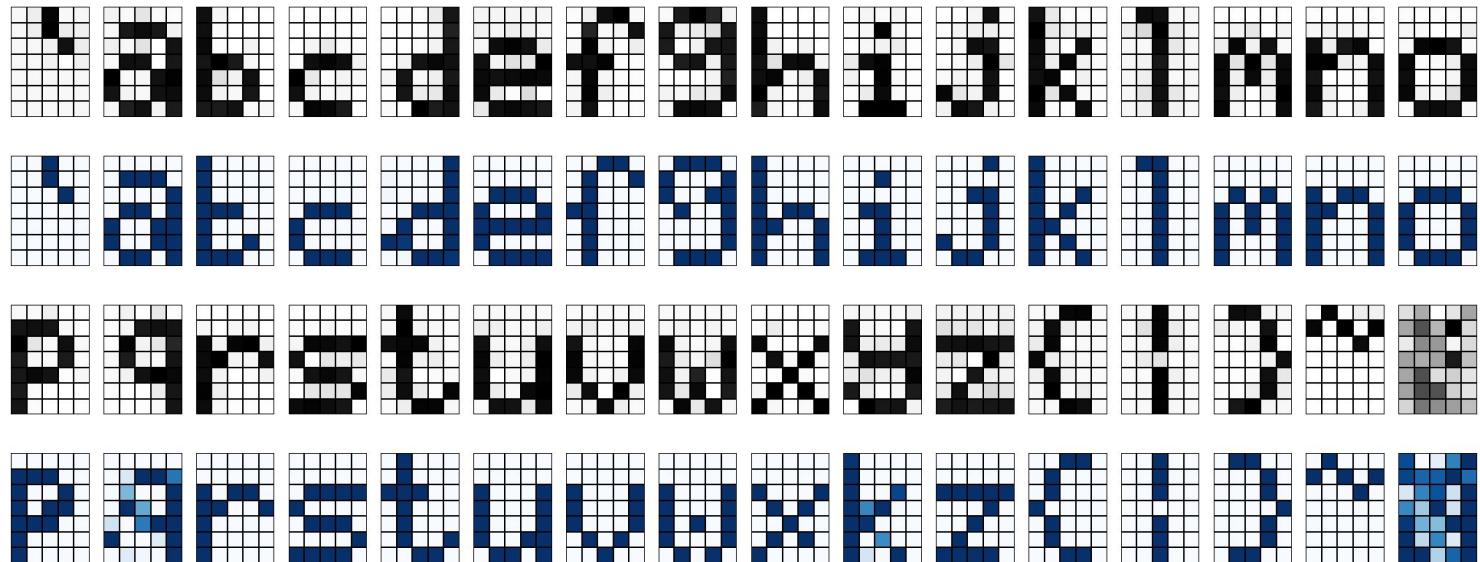
Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None

Arquitectura: [60, 50, 30, 10, 5]

Dataset - Ruido 0.1

Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None

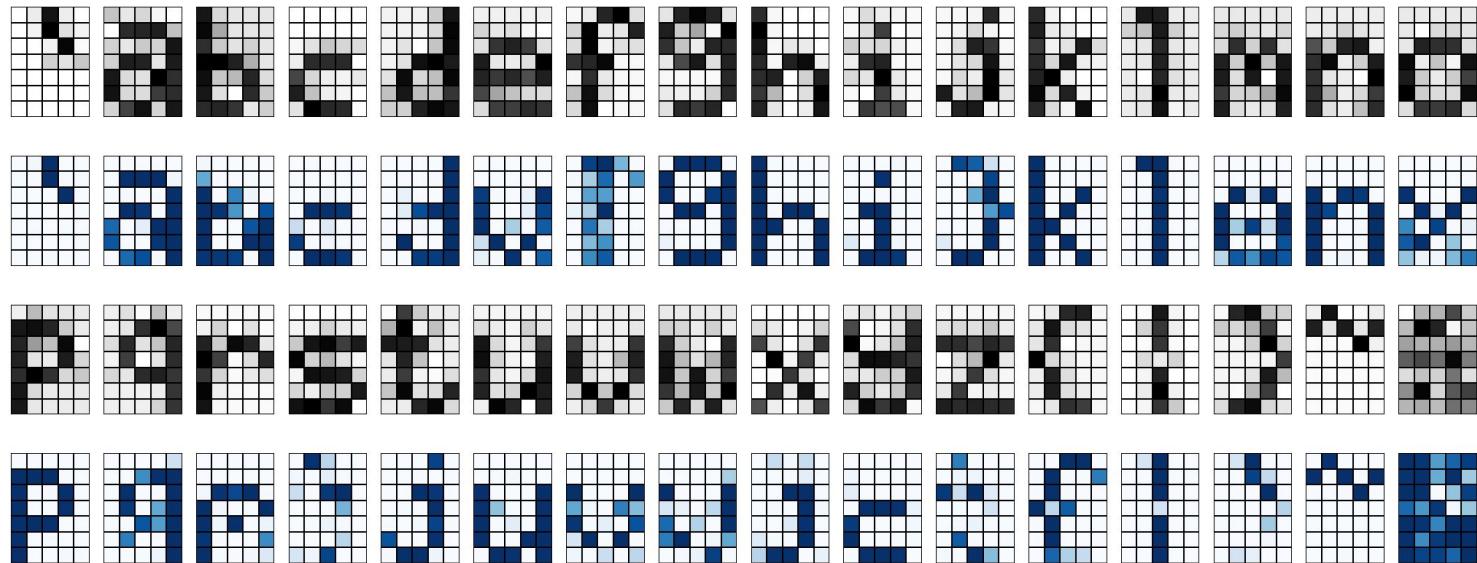


Arquitectura: [60, 50, 40, 30, 20, 10, 5]

Dataset - Ruido 0.3

Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None



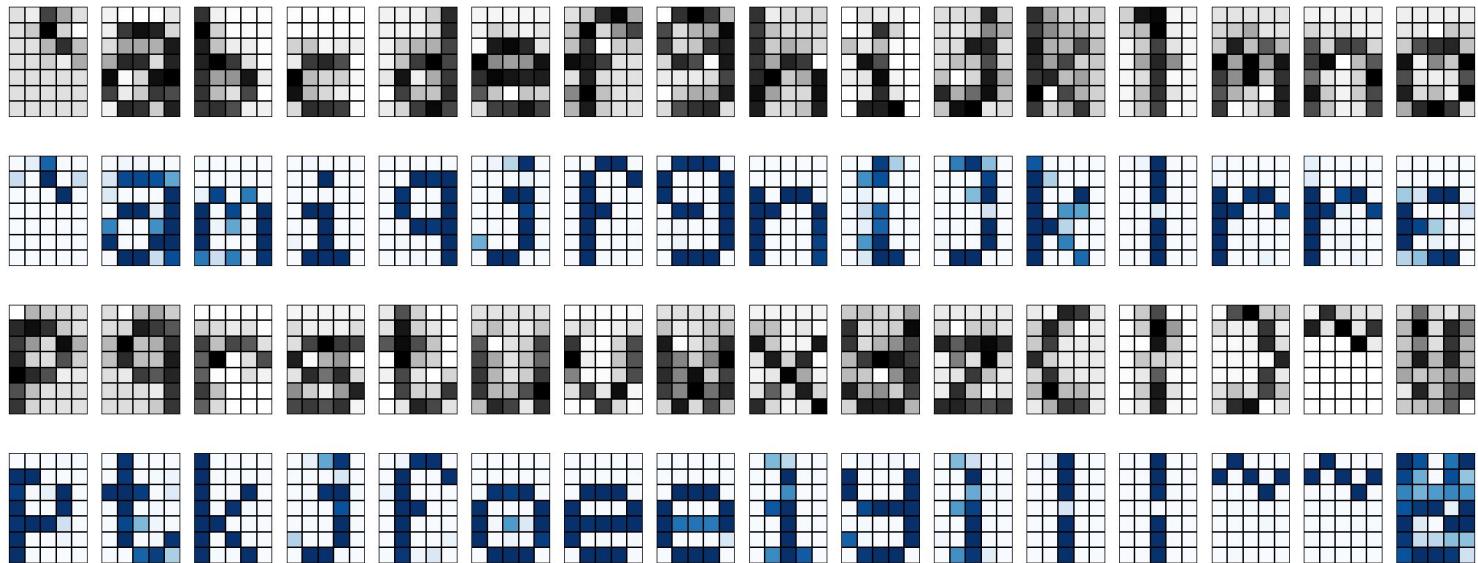
Arquitectura: [60, 50, 30, 10, 5]



Dataset - Ruido 0.5

Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise = None

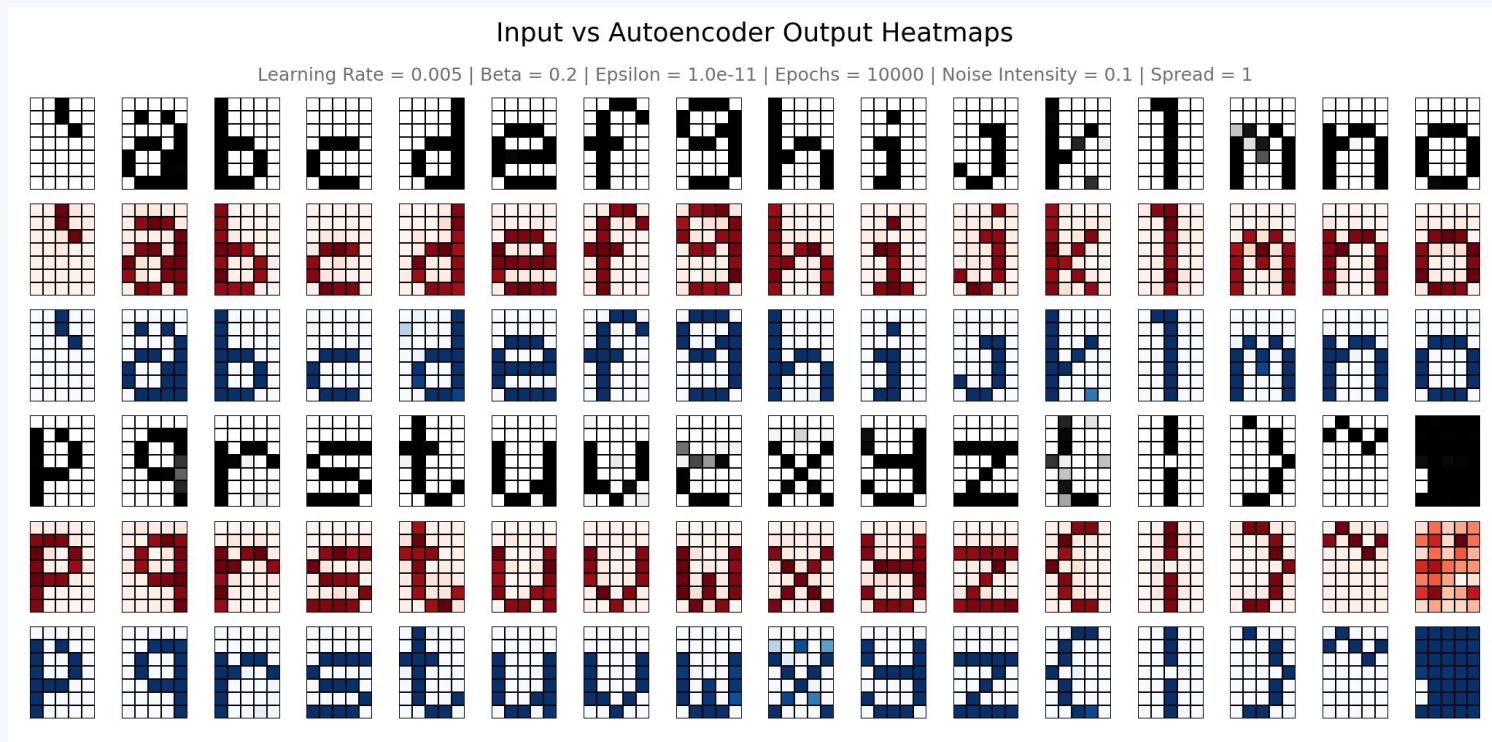


Arquitectura: [60, 50, 30, 10, 5]



Ruido 0.1

- Aprendido
(paso previo)
- Entrada con
ruido
- Interpretación

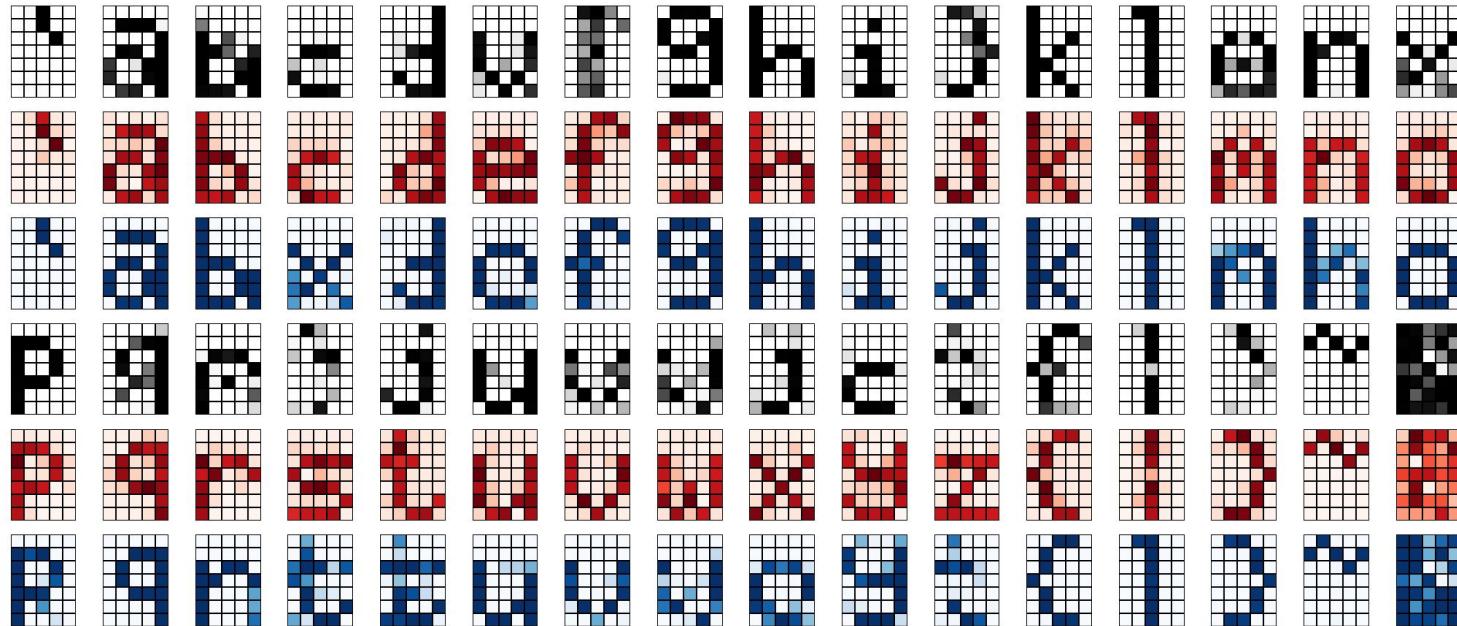




Ruido 0.3

Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise Intensity = 0.3 | Spread = 1



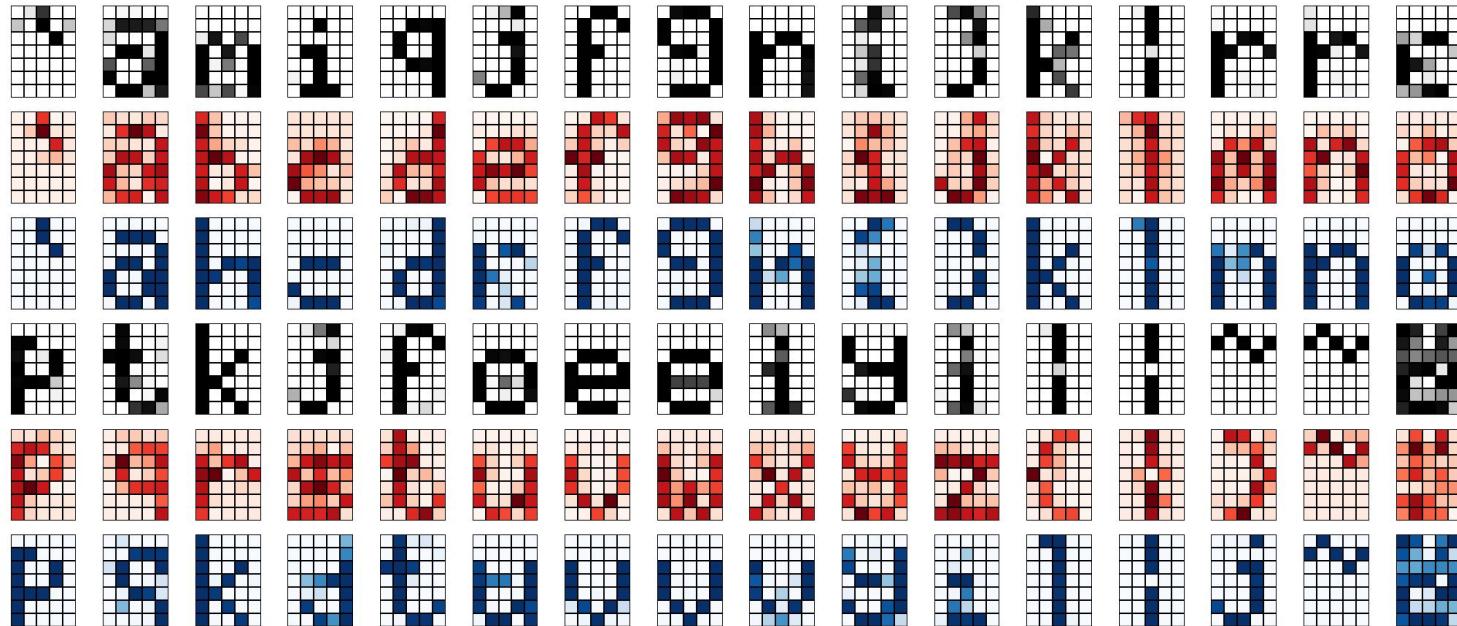
Arquitectura: [60, 50, 30, 10, 5]



Ruido 0.5

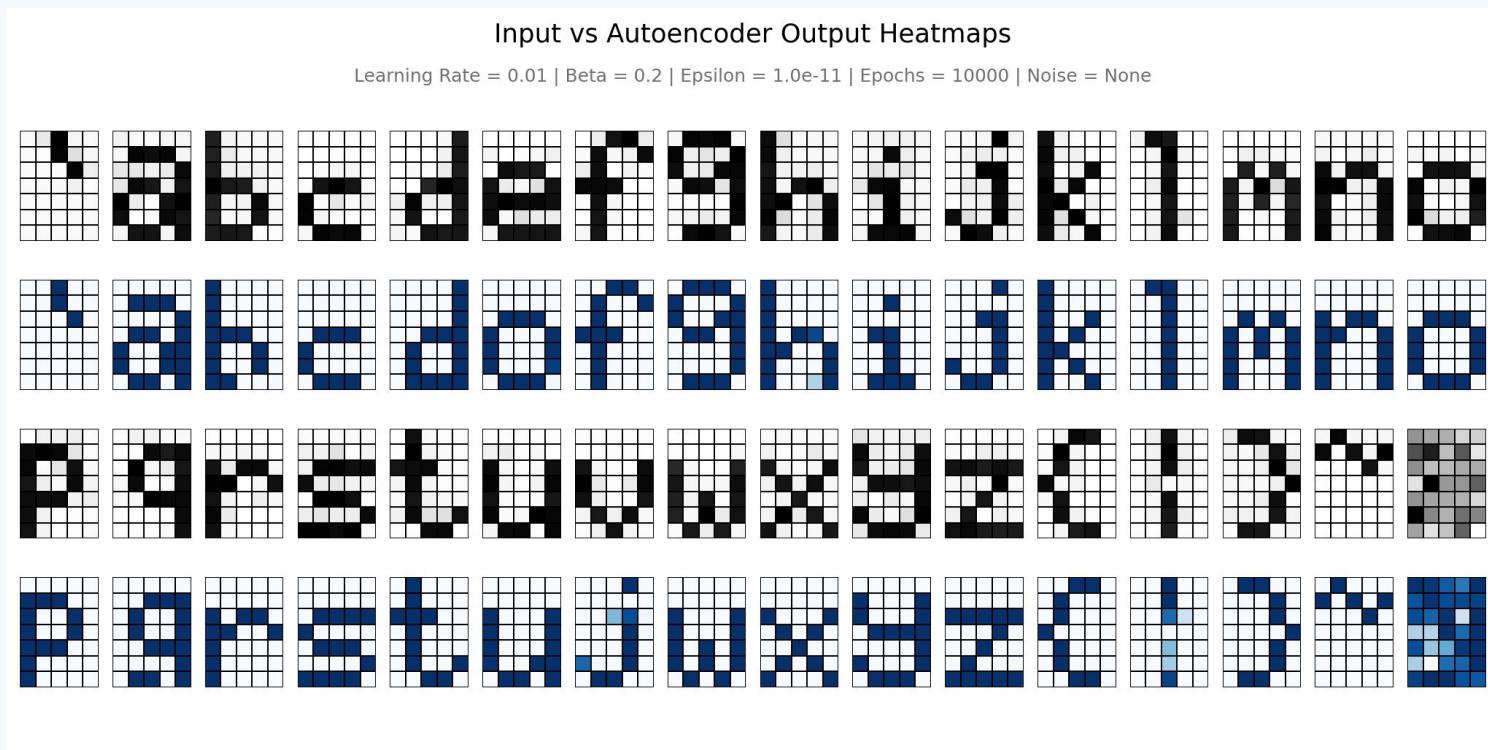
Input vs Autoencoder Output Heatmaps

Learning Rate = 0.005 | Beta = 0.2 | Epsilon = 1.0e-11 | Epochs = 10000 | Noise Intensity = 0.5 | Spread = 1



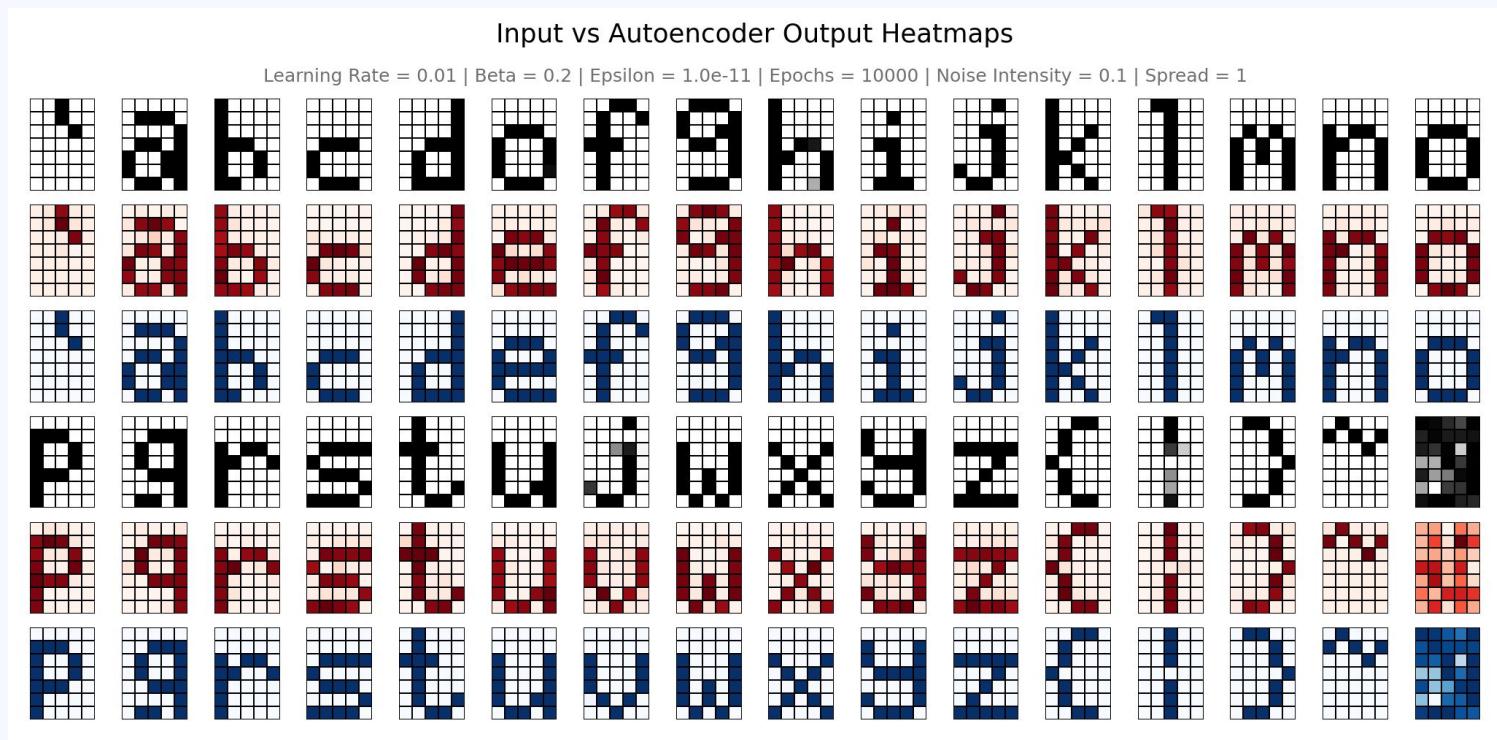
Arquitectura: [60, 50, 30, 10, 5]

Dataset - Ruido 0.1 - LR = 0.01



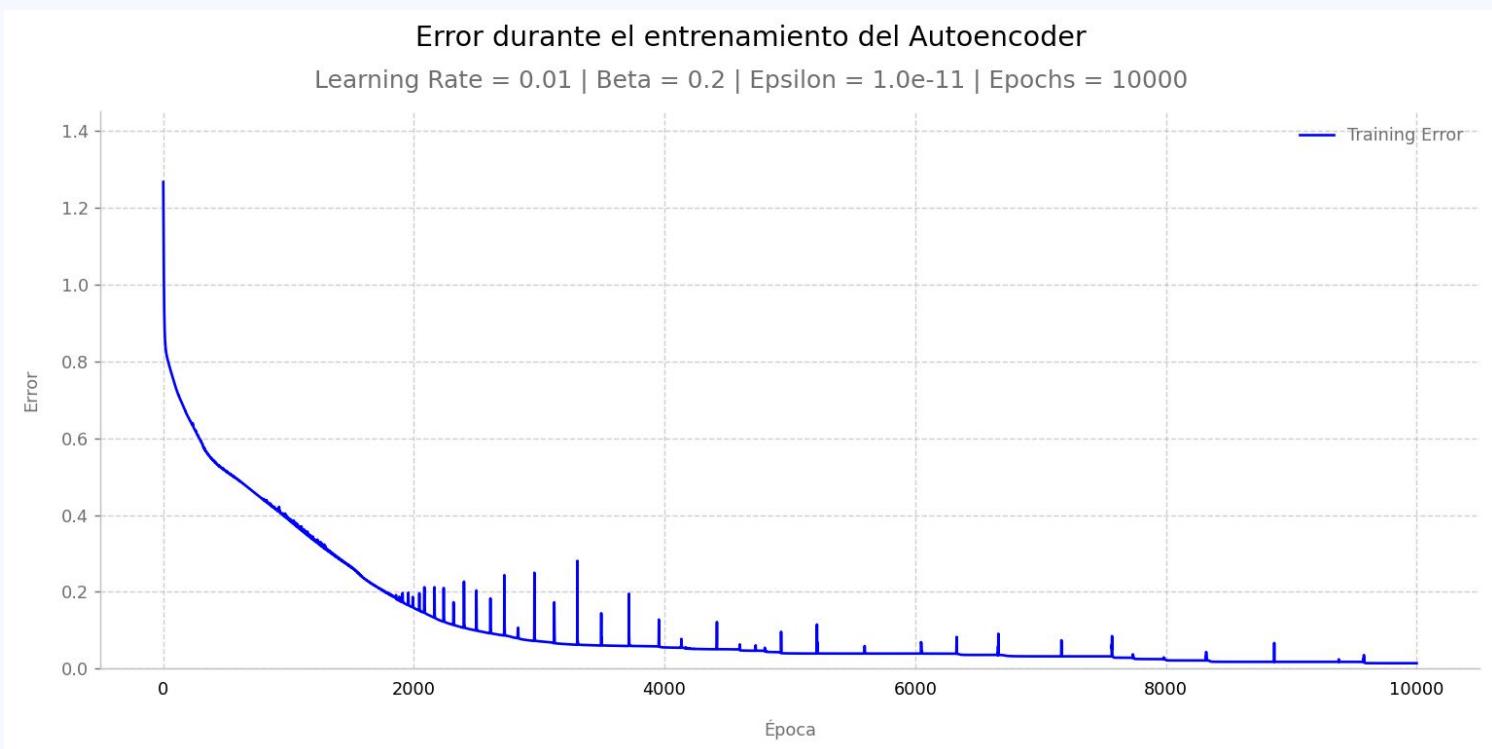
Ruido 0.1 - LR = 0.01

- Aprendido
(paso previo)
- Entrada con
ruido
- Interpretación



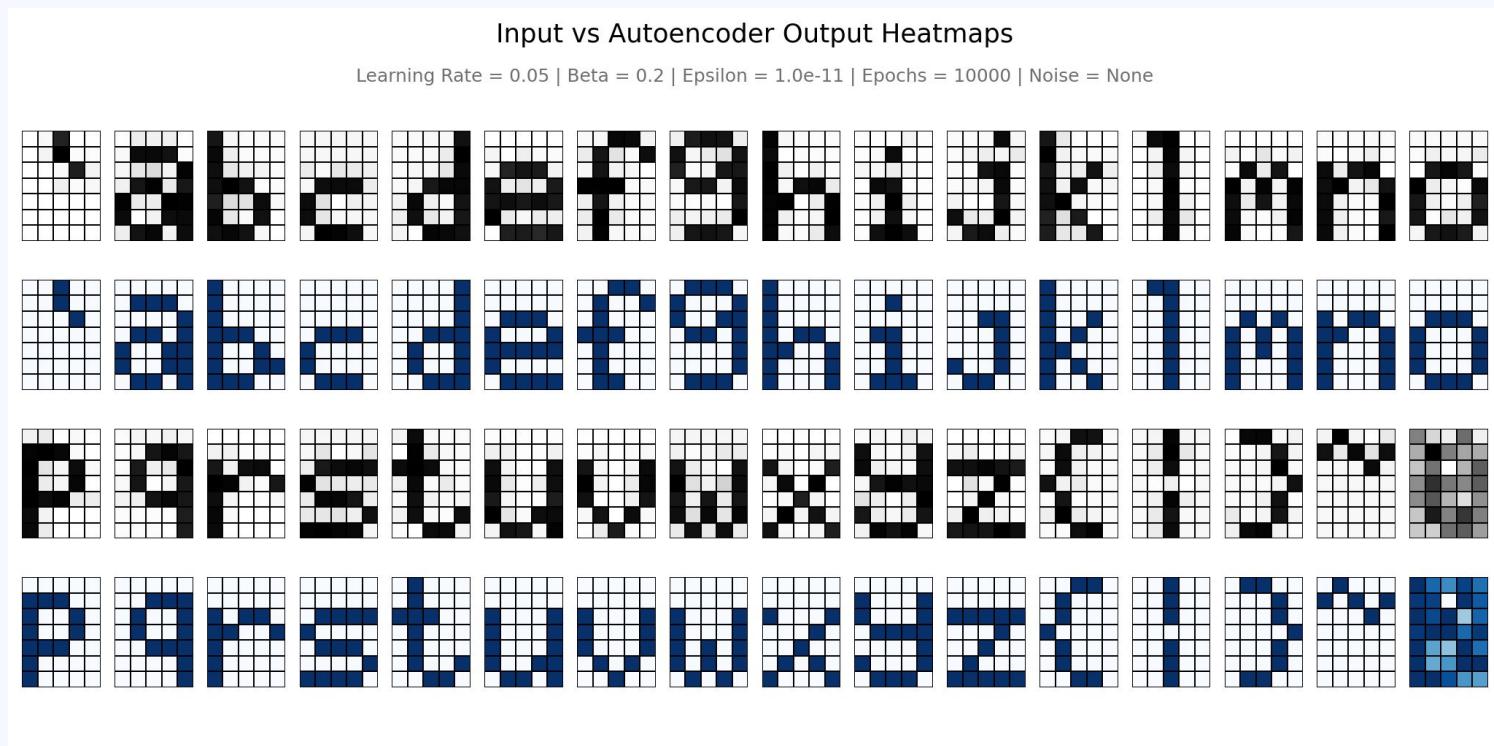
Arquitectura: [60, 50, 30, 10, 5]

Ruido 0.1 - LR = 0.01



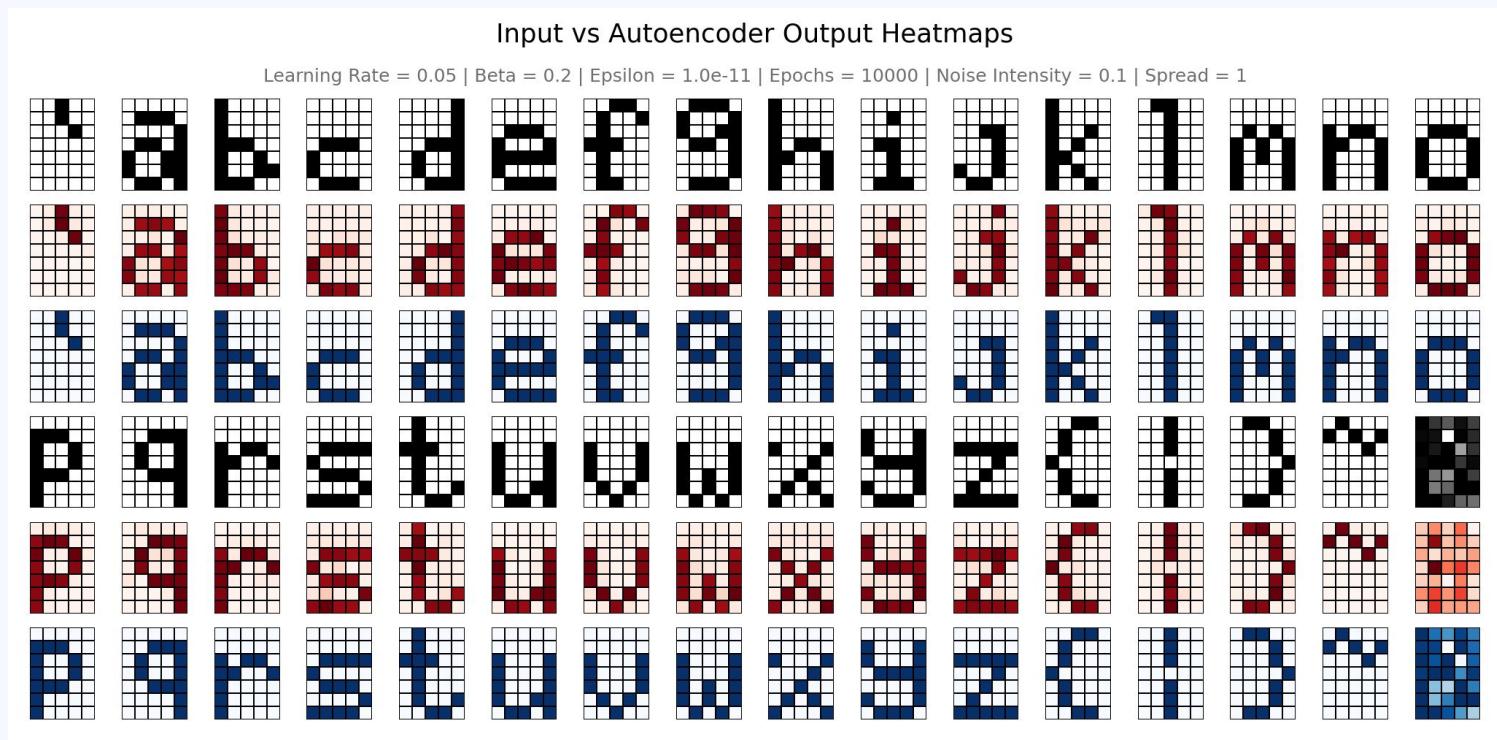
Arquitectura: [60, 50, 30, 10, 5]

Dataset - Ruido 0.1 - LR = 0.05



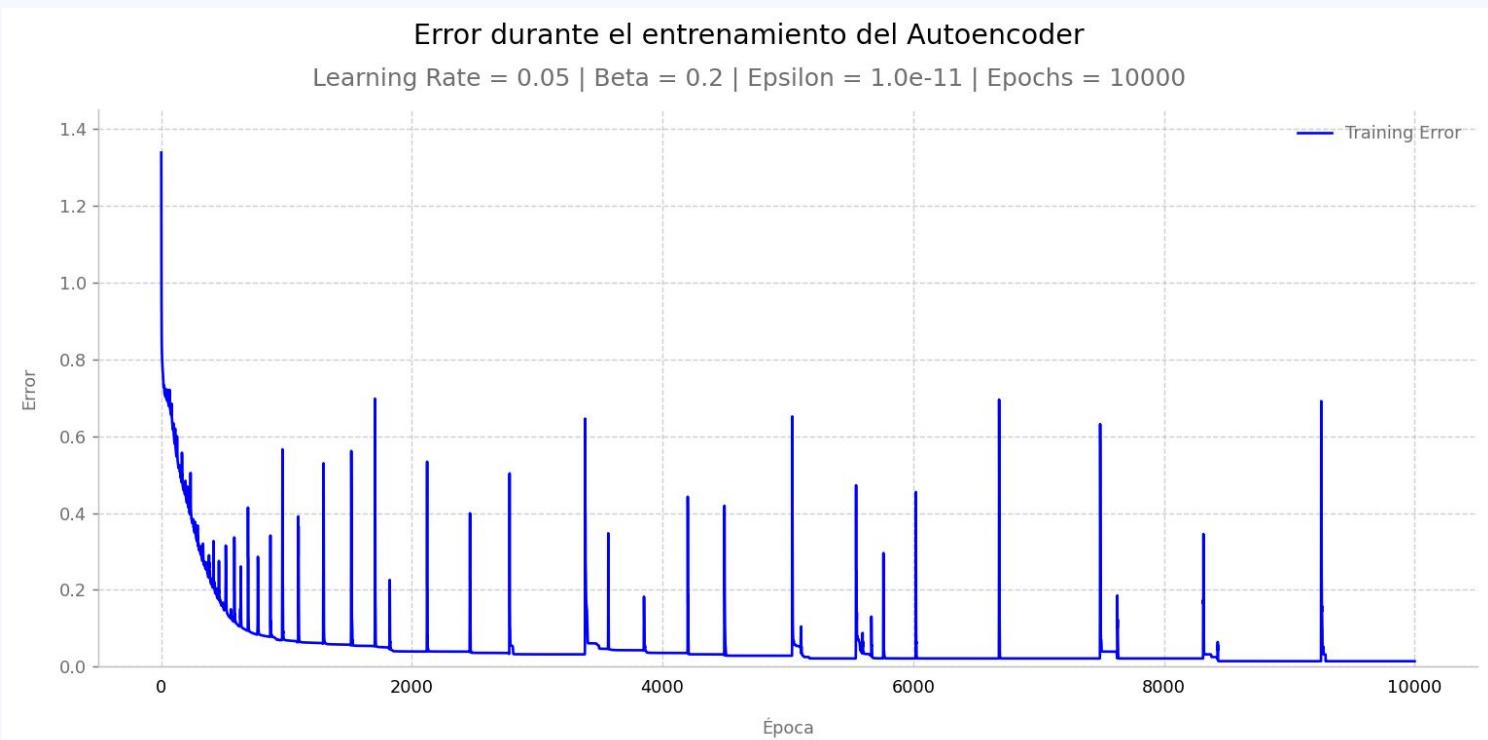
Arquitectura: [60, 50, 30, 10, 5]

Ruido 0.1 - LR = 0.05



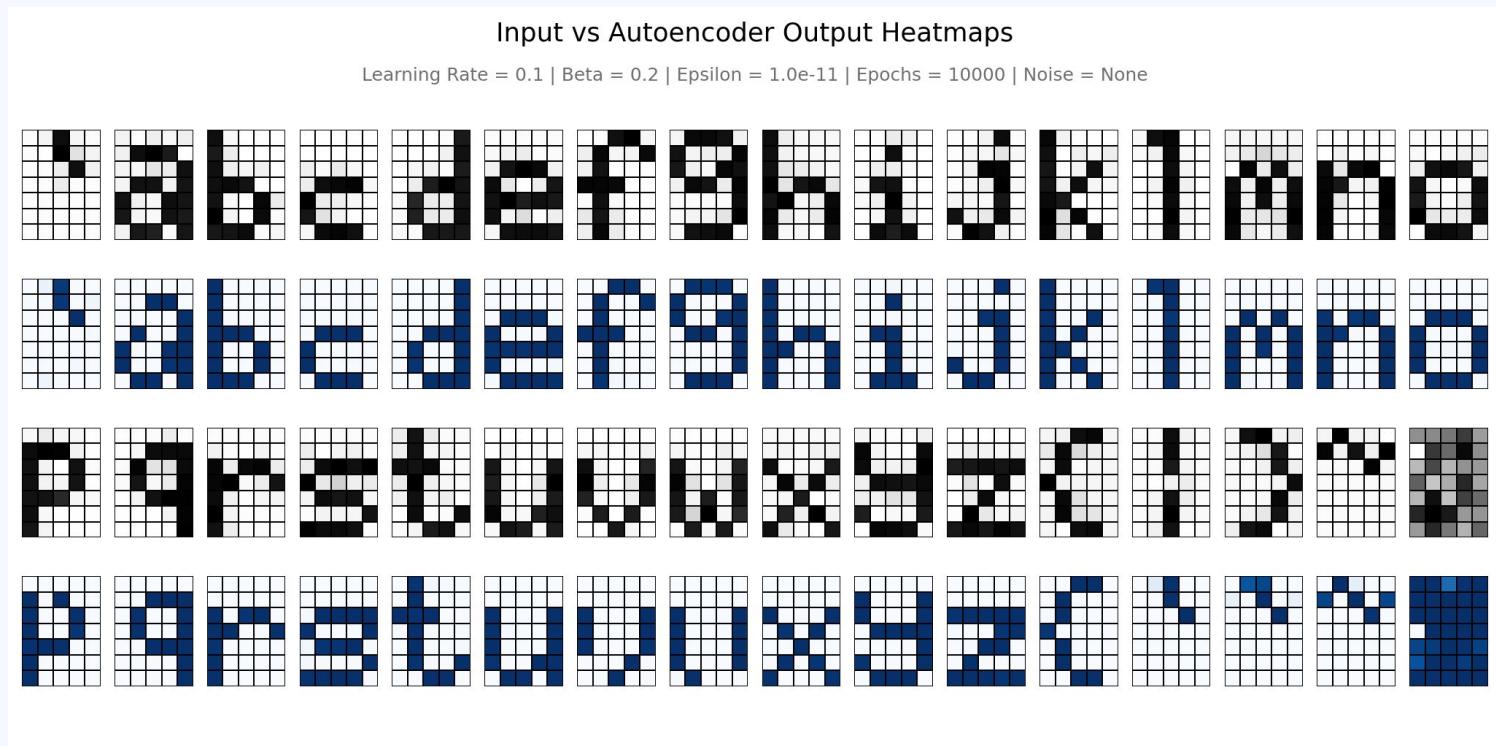
Arquitectura: [60, 50, 30, 10, 5]

Ruido 0.1 - LR = 0.05



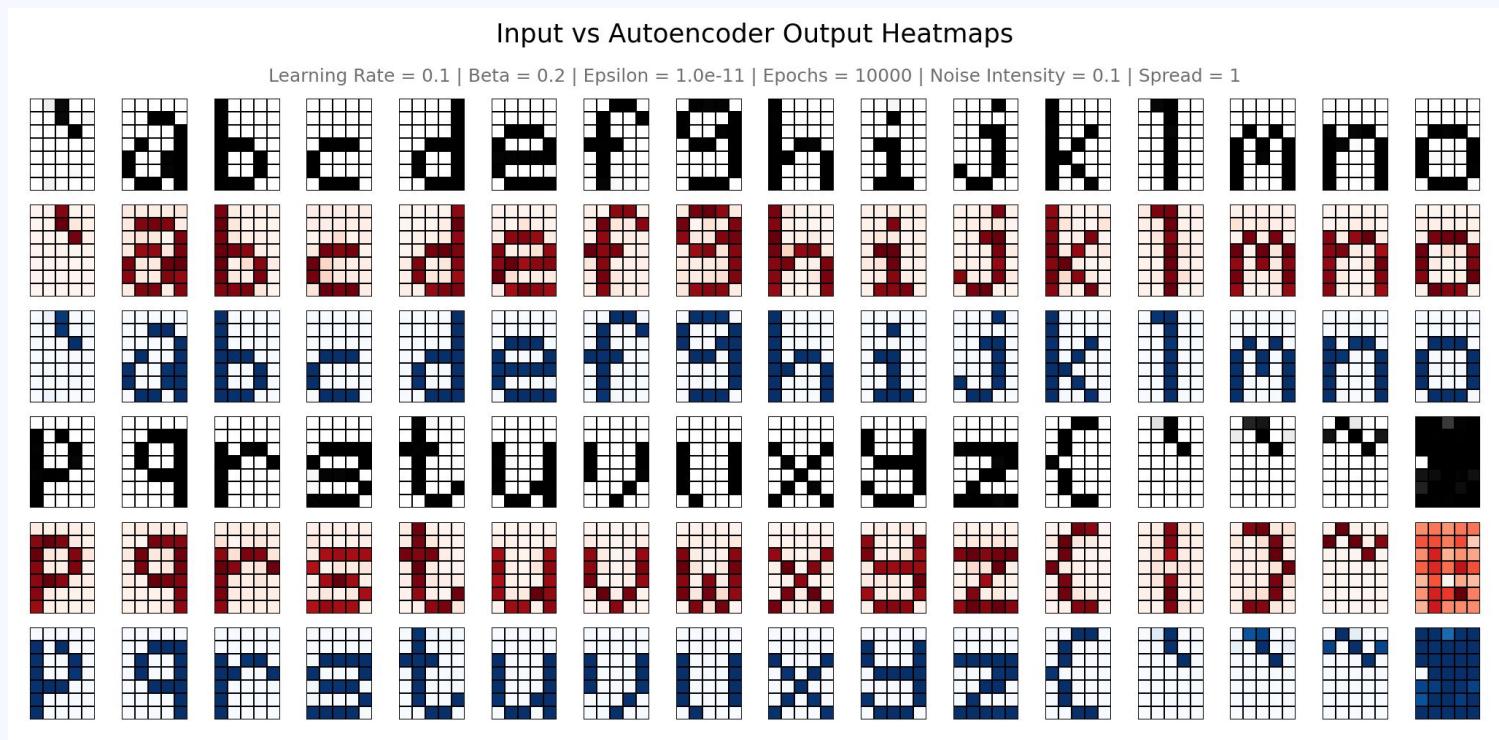
Arquitectura: [60, 50, 30, 10, 5]

Dataset - Ruido 0.1 - LR = 0.1



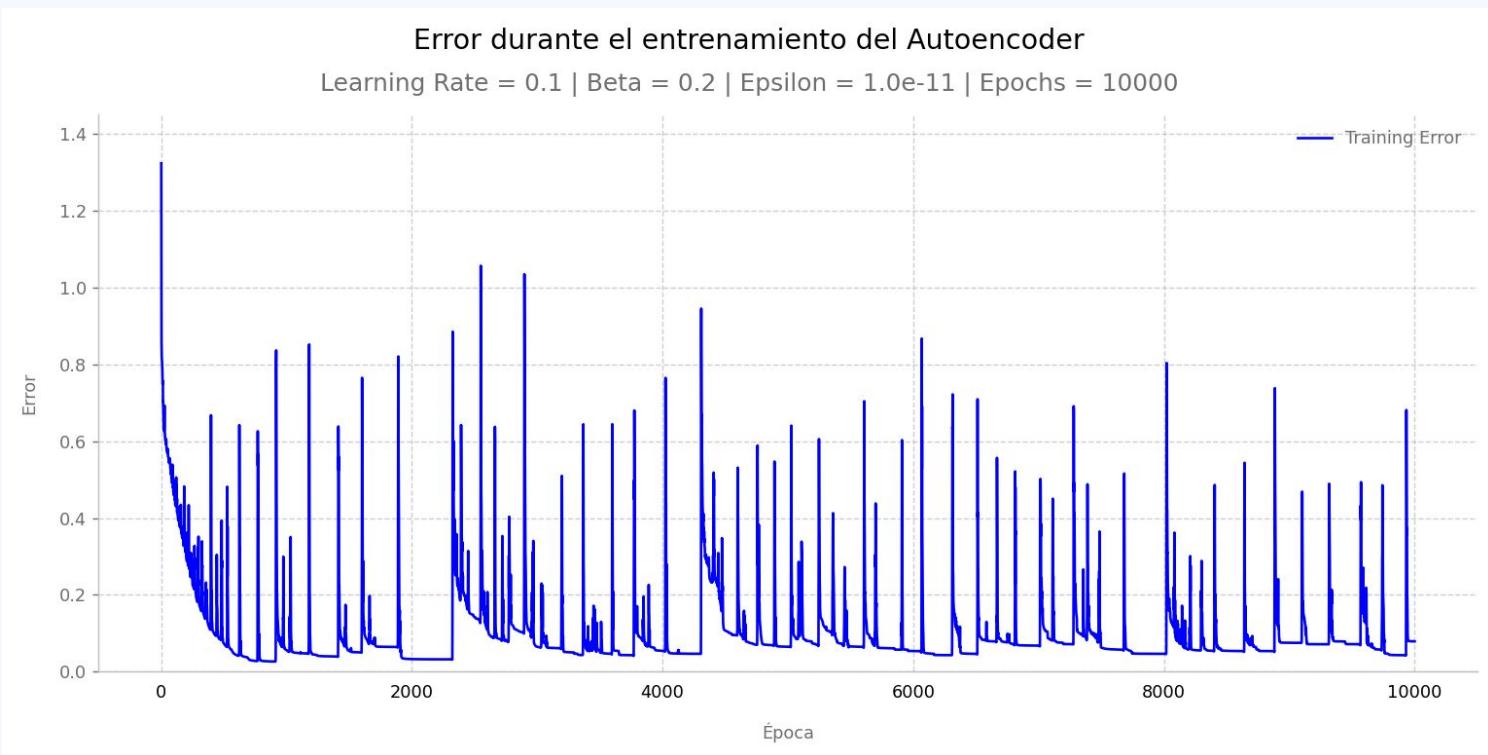
Arquitectura: [60, 50, 30, 10, 5]

Ruido 0.1 - LR = 0.1



Arquitectura: [60, 50, 30, 10, 5]

Ruido 0.1 - LR = 0.1



Arquitectura: [60, 50, 30, 10, 5]

Conclusiones

- Los autoencoders pueden detectar y filtrar muy bien el ruido
- Cambiar la arquitectura, learning rate y cantidad de épocas altera significativamente el resultado y tiempos de ejecución
- ¡Cuidado con el error!

03

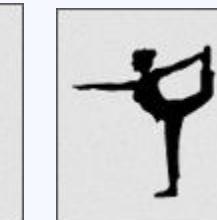
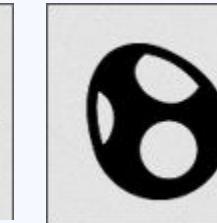
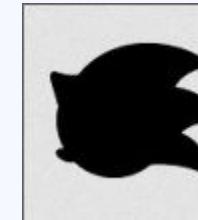
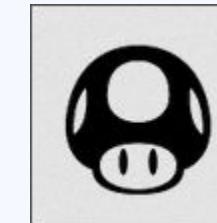
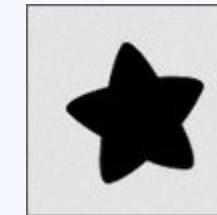
Autoencoder

Variacional



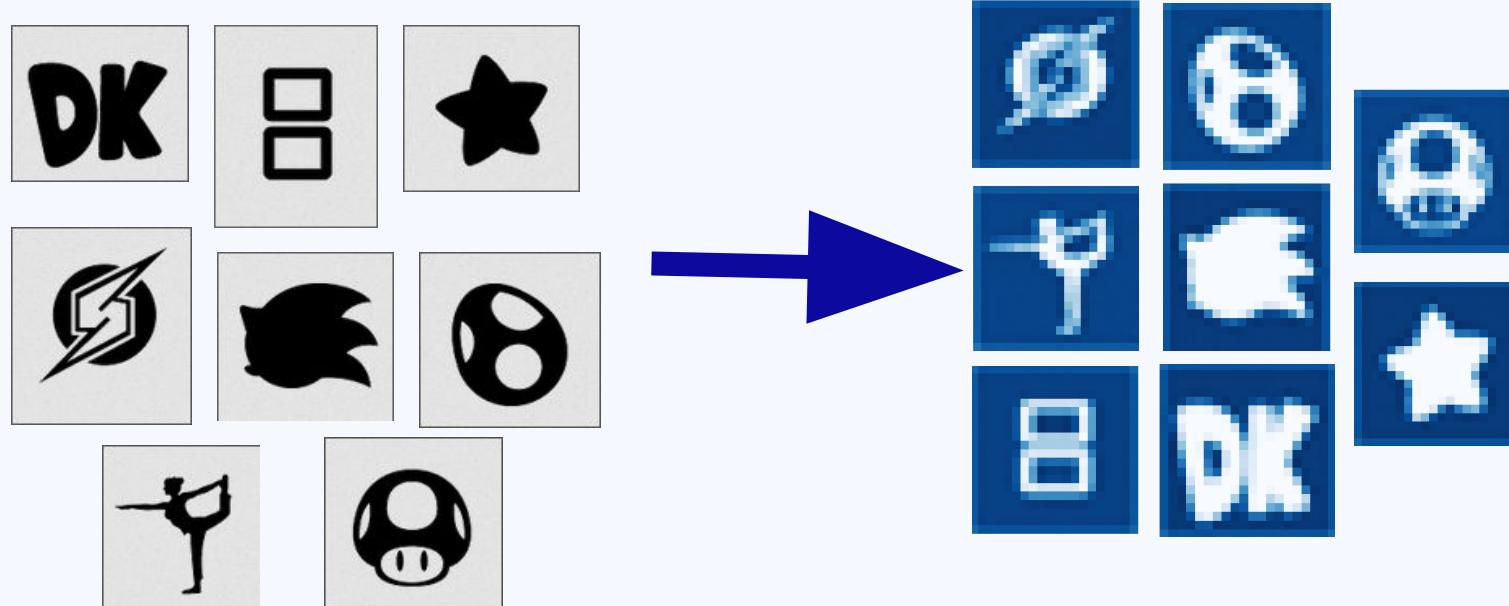
Conjunto de datos nuevo

Para este conjunto de datos elegimos los logotipos de alguna de las franquicias de Nintendo



Conjunto de datos nuevo

Para poder entrenar a la red, procesamos las imágenes para que tengan un mismo tamaño y resolución, y los cambiamos a una escala de azules



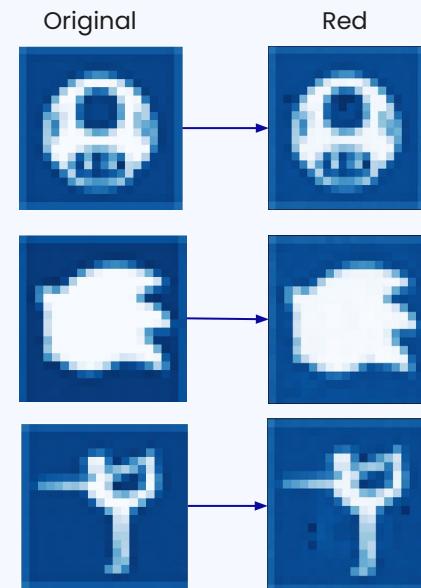
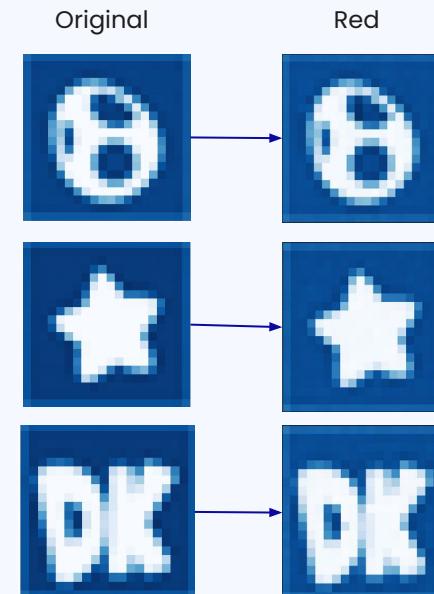
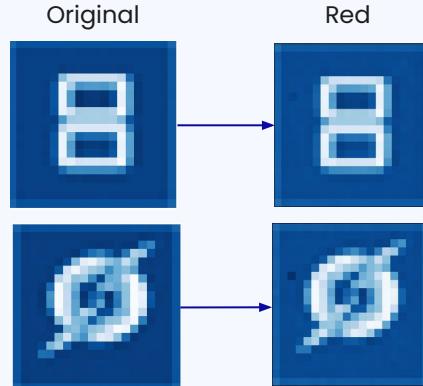
Algunos Resultados

Arquitectura: **400-300-200-100-50-2-50-100-200-300-400**

Learning Rate: 0.005

Epochs: 1500

Optimizer: Adam



Generación de Nuevos Logos

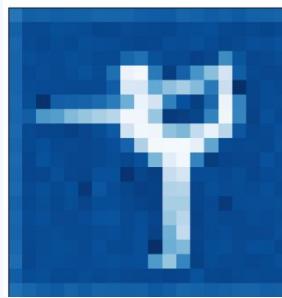
Para poder generar nuevos logos, codificamos las imágenes y generamos sus valores en el espacio latente.

```
latent_space_1 = autoencoder.encode(img1)  
latent_space_2 = autoencoder.encode(img2)
```

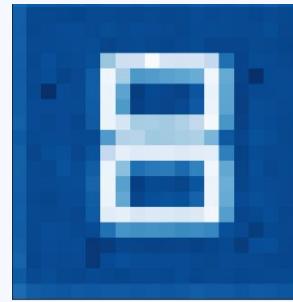
Luego, decodificamos la imagen resultante del promedio de sus valores:

```
mean_latent_space = 0.5 * (latent_space_1 + latent_space_2)  
new_img = autoencoder.decode(mean_latent_space)
```

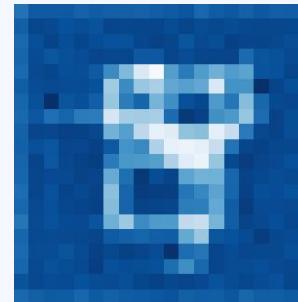
Generación de Nuevos Logos



+



=



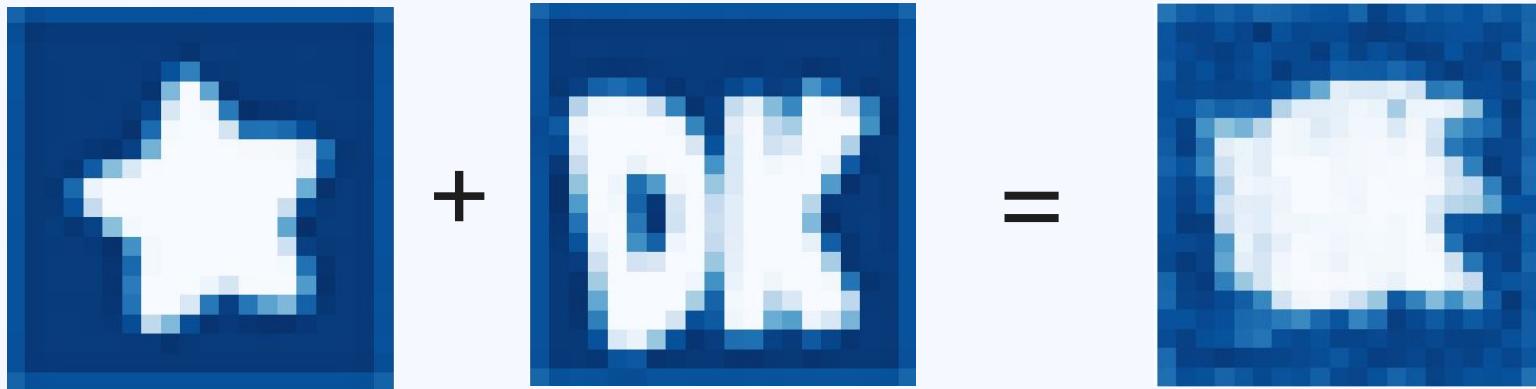
+



=



Generación de Nuevos Logos



Conclusiones

- Los VAEs permiten generar nuevas muestras y realizar interpolaciones más suaves sobre el espacio latente.
- La representación del espacio latente en un autoencoder no variacional no tiene una estructura probabilística, lo que limita la capacidad de generar nuevas muestras.

Muchas gracias!
